

Towards Energy-Aware Design Choices in Event-Driven Microservices

Vijay
Roll No. 2023111007
IIIT Hyderabad
Hyderabad, India

Abstract—As enterprise software shifts to cloud-native microservices, architectural decisions like communication topology, message broker selection, and data serialization are usually optimized for speed and scalability. However, these choices directly impact hardware resource consumption (CPU, memory, disk I/O, and network bandwidth), which translates to physical energy use and carbon emissions. This paper evaluates three foundational pillars of event-driven architectures (EDA) through the lens of resource-use proxies for the Software Carbon Intensity (SCI) specification. Through controlled simulation, we indicate that choosing parallel fan-out topologies over sequential chains reduces simulated event latency by 66%, replacing JSON with a packed binary surrogate reduces payload size by 66% and parsing CPU overhead by nearly 90%, and selecting the appropriate message broker based on throughput alters the baseline CPU footprint. These experiments suggest that lower-latency topologies and more compact binary encodings can reduce resource-use proxies that may translate into lower energy consumption under an explicitly stated power model, providing guidance for sustainable software architecture.

Index Terms—Green Software Engineering, Microservices, Event-Driven Architecture, Energy Efficiency, Software Carbon Intensity

I. Introduction

As software systems increasingly transition toward cloud-native microservices and event-driven architectures (EDA), decisions regarding service topology, message broker selection, and data serialization formats are traditionally guided by scalability, resilience, and responsiveness. However, these structural choices profoundly influence CPU cycles, network traffic, memory pressure, and message broker overhead—proxies that correlate with physical hardware energy consumption.

For organizations deploying cloud-native distributed systems at scale, minor inefficiencies in per-event processing can compound into significant environmental and infrastructure costs. The motivation for this research is to evaluate software design patterns for microservices through a rigorous sustainability lens, recognizing that macro-level architectural choices often exert a significant long-term impact on resource utilization.

This project proposes a simulation study to determine whether common event-driven design decisions can be systematically ranked by their resource consumption per business event. By quantifying the tradeoffs inherent in topology, broker selection, and serialization using proxy

metrics, we provide a theoretical framework for estimating the operational energy component of the Software Carbon Intensity (SCI) specification in practice. To promote reproducibility and green software practices, all benchmarking scripts, mathematical models, and datasets developed for this study are fully open-source and publicly available at: <https://github.com/VA24d/SE-bonus>.

II. Literature Review

To position this empirical simulation study, recent academic research on green software engineering provides a concrete foundation for evaluating architectural energy footprints.

A. Green Software Engineering and SCI

The Software Carbon Intensity (SCI) specification, officially codified as ISO/IEC 21031:2024, provides a standard methodology for scoring the carbon footprint of software applications. SCI measures carbon emissions per functional unit, establishing an intensity rate using the equation:

$$SCI = \frac{(E \times I) + M}{R} \quad (1)$$

where E is operational energy, I is the grid carbon intensity, M is the embodied carbon of the hardware, and R represents the functional unit (e.g., events processed). Reducing the CPU and network overhead directly lowers the operational energy (E) component of the SCI score [1], [2].

B. Event-Driven Microservices Topologies

Microservice architectures differ radically in how services interact. Empirical studies consistently find that asynchronous choreography yields lower end-to-end latency and CPU utilization than strict synchronous orchestration. For instance, quantitative measurements by Ristova et al. evaluating six canonical service-dependency topologies observed that parallel fan-out structures (highly asynchronous) were the most energy-efficient, while sequential chain designs consumed significantly more energy (up to nearly 39 kJ at scale) [4]. Orchestrated call chains incur idle CPU blocking overhead, whereas event-driven choreography leverages concurrency to improve resource utilization.

C. Message Broker Performance

Message brokers exhibit different resource profiles based on their architecture. Benchmark reports indicate that Kafka achieves much higher throughput than RabbitMQ, but with higher baseline resource use. Kafka can sustain roughly 76,000–100,000+ messages per second compared to 7,000–18,000 msg/s for RabbitMQ [5]. Consequently, Kafka’s theoretical CPU usage per message is lower at massive scales, largely due to its utilization of the sendfile zero-copy system call and sequential disk I/O, which bypasses application-level CPU context switching. Conversely, RabbitMQ is highly resource-efficient for low-volume scenarios due to its in-memory nature, allowing it to use nearly 29% less memory than Kafka for lightweight loads.

D. Serialization Formats

The choice of data format directly affects CPU and network proxy metrics. A recent ICSoft 2025 study by Shatnawi et al. refactored web services from JSON to Protocol Buffers (Protobuf), reporting a 60–80% reduction in payload size, an 80% faster response time, 17% lower CPU utilization, and an 18–33% reduction in energy consumption [6]. This quantifies the “JSON tax” consisting of extra CPU cycles for text parsing and network energy for larger payloads, suggesting that binary schemas provide a clear resource win in high-throughput environments [3].

III. Methodology and Experimental Design

To provide reproducible proxy metric estimates, we conducted three controlled simulations mimicking realistic microservice workloads. All software simulations were developed in Python 3.10 and executed on an ARM64-based macOS environment (Apple Silicon). We utilized standard libraries such as `time.perf_counter()` for high-resolution CPU time tracking and `asyncio` for parallel non-blocking I/O execution.

A. Serialization Energy Benchmarking

We designed a Python-based workload simulator that generated 1,000,000 standardized business events. Each event mimicked an e-commerce order containing randomly generated data structures (a UUID string, an integer user ID, a short status string, and a float numerical amount). We benchmarked three distinct serialization formats: standard JSON (text-based), MessagePack (a binary, schema-less format), and a tightly packed Python Struct (serving as a packed binary surrogate to simulate the strict binary compactness and fixed-width fields of Protocol Buffers).

For each format, we flushed the system cache and executed the serialization and deserialization functions 1,000,000 times sequentially. Metrics gathered included the total serialized payload footprint in memory (measured in MB), execution time (measured in ms), and actual physical energy consumed (kWh) alongside carbon

emissions ($kgCO_2eq$). This was achieved by wrapping the runtime execution in the CodeCarbon emissions tracker, which actively polls Intel RAPL/Apple Silicon hardware telemetry to record true processor power draw rather than relying solely on CPU time proxies.

B. Physical Network Topology Simulation

To evaluate the impact of microservice orchestration versus choreography under true network contention, we deployed three physical Python ThreadingHTTPServer instances on discrete local ports representing independent microservices. We assumed both topologies process the exact same core business semantics (processing a batch of 50 events across the 3 services). We defined two microservice communication patterns representing idealized bounds: 1. Sequential Chain (Orchestration): Service A explicitly waits for Service B to return, which in turn waits for Service C. Thread execution blocks sequentially until the entire chain is resolved. 2. Parallel Fan-Out (Choreography): An event is emitted to a central broker construct, triggering Services A, B, and C to independently consume and process the event simultaneously without blocking the parent caller.

We attributed a fixed 50ms I/O-bound delay per service processing step, and artificially injected 10ms of network latency per HTTP request to simulate physical distributed I/O. We measured the total elapsed end-to-end latency to process the batch. Higher latency in the sequential chain correlates with prolonged thread blocking and idle CPU wait states, which decrease processing density per watt.

C. Broker Overhead Mathematical Modeling

Because executing a massive-scale physical infrastructure benchmark requires enterprise server clusters, we designed a robust mathematical simulation model. This model predicts message broker CPU utilization scaling based on established empirical data gathered from industry benchmarking [5].

- RabbitMQ Model: Configured with a low baseline CPU footprint (B_{RMQ}) typical of the Erlang VM, but modeled with an exponentially scaling overhead to replicate the known limits of its memory pressure and state tracking under high-throughput loads. The theoretical CPU utilization U as a function of message rate N is defined as:

$$U_{RMQ}(N) = 2.0 + 4.5 \cdot \left(\frac{N}{1000}\right)^{1.5}$$

- Kafka Model: Configured with a significantly higher baseline CPU footprint (B_{Kafka}) accounting for the JVM and continuous disk polling, but scaling linearly. This reflects its architectural advantage of log batching and zero-copy file transfers. The utilization function is defined as:

$$U_{Kafka}(N) = 12.0 + 1.5 \cdot \left(\frac{N}{1000}\right)$$

We computationally simulated throughput ranges from 1,000 to 50,000 messages per second, plotting the theoret-

ical efficiency cross-over point where scale overcomes the baseline power footprint.

IV. Results and Empirical Analysis

The following subsections present the concrete findings from our simulations, visually demonstrating the profound differences in resource-use proxies between architectural choices.

A. Serialization Impact

As seen in Figure 1, migrating from a text-based format like JSON to a binary format reduces both the network payload and the computational overhead required for parsing.




Fig. 1. Comparison of payload size and CPU processing time across 1,000,000 business events using JSON, MsgPack, and a packed binary surrogate.

The 1,000,000 JSON events consumed 67.71 MB of bandwidth and required 1362.57 ms to process, consuming 0.000031 kWh of physical energy (emitting 0.000022 kg CO₂eq). MessagePack served as an effective middle-ground, dropping the payload to 49.59 MB and execution time to 468.34 ms (0.000012 kWh) without requiring a strict schema. However, the tightly packed binary surrogate yielded the highest physical efficiency, consuming only 22.89 MB, taking 116.86 ms, and drawing just 0.000004 kWh. This represents an 85.5% reduction in actual physical energy compared to JSON. When multiplied across millions of daily microservice interactions, eliminating the "JSON tax" directly lowers the physical metrics associated with operational energy consumption.

B. Topology Latency Impact

Figure 2 illustrates the total latency to process 50 business events across three independent services depending on the chosen communication topology.

The Sequential Chain topology, representing an orchestrated request path with absolute data dependency coupling across HTTP, resulted in 10.15 seconds of total latency due to continuous thread blocking and idle wait times across the 10ms network boundary. By refactoring




Fig. 2. End-to-end latency comparison between simulated sequential chain requests and parallel fan-out event-driven processing.

to a Parallel Fan-Out (choreographed) topology where services execute independently via concurrent HTTP requests, the exact same computational work was completed in 3.60 seconds.

This 64.5% reduction in latency aligns with the literature findings [4]: decoupling services and leveraging parallel event consumption minimizes idle CPU blocking overhead. By freeing up threads faster and improving CPU utilization density, the hardware processes more functional units per watt of energy.

C. Message Broker Efficiency Scaling

The mathematical simulation of message broker efficiency reveals a clear theoretical cross-over point where the architectural design of the broker becomes the dominant factor in CPU proxy metrics.




Fig. 3. Simulated CPU utilization comparison of RabbitMQ vs. Apache Kafka across increasing message throughput rates.

As depicted in Figure 3, RabbitMQ is highly resource-efficient at low throughputs (e.g., 1,000 msg/s), consuming only 6.5% CPU compared to Kafka's 13.5% baseline.

However, as load increases to 5,000 msg/s, RabbitMQ’s theoretical CPU utilization spikes to 52.3% due to in-memory state tracking overhead, whereas Kafka’s batching mechanisms keep its modeled CPU usage at an efficient 19.5%.

Beyond 10,000 msg/s, the RabbitMQ model becomes fully saturated (overload zone), while the Kafka model gracefully scales up to 50,000 msg/s without exhausting the CPU constraint. Under the stated model assumptions, choosing RabbitMQ may optimize resource use for lightweight, low-volume systems to save baseline power, whereas Kafka is indicated for enterprise-scale systems to minimize resource utilization per message.

V. Future Work

While this simulation study effectively isolates the proxy metrics associated with microservice design patterns, our future work aims to physically deploy these architectures to calculate empirical energy consumption. The planned experimental scope focuses on the following primary objectives:

- Physical Microservice Testbed: Deploying a controlled testbed (comprising 4 to 6 discrete services) via Docker containers onto a standardized, bare-metal infrastructure rather than simulating workloads on a local machine.
- Hardware Energy Telemetry: Instrumenting the physical infrastructure to record actual power draw (Watts) using dedicated energy telemetry tools, capturing true physical data instead of relying on CPU utilization and memory payloads as proxies.
- SCI Metric Mapping: Mapping the empirical energy consumed per functional unit (e.g., per 1,000 processed events) directly to the formal Software Carbon Intensity (SCI) specification, factoring in the regional carbon intensity of the electricity grid and embodied hardware emissions.

This will transition our findings from a theoretical resource model into a fully realized, empirical benchmark for green software engineering.

VI. Conclusion

By simulating the proxy metrics tied to core architectural decisions, this study suggests that sustainability is influenced by system design long before code-level optimizations occur. Transitioning from synchronous JSON chains to parallelized, binary-encoded, Kafka-backed streams can reduce resource-use proxies (CPU time, thread blocking, payload footprint) that correlate with hardware power models. This research provides a theoretical framework to empower software engineers to integrate sustainability as a core non-functional requirement, utilizing these guidelines to optimize the operational energy components of modern cloud-native applications.

References

- [1] R. Mahmoud and A. Qusef, “A Systematic Literature Review on Green Software Engineering,” *IEEE Access*, vol. 10, pp. 8652–8664, 2022.
- [2] Green Software Foundation, “Software Carbon Intensity (SCI) Specification (ISO/IEC 21031:2024),” 2024. [Online]. Available: <https://greensoftware.foundation/sci>. [Accessed: 30-Apr-2026].
- [3] S. Hasan et al., “Performance and Energy Efficiency of Serialization Formats in Microservices,” *Journal of Cloud Computing*, vol. 12, no. 1, p. 34, 2023.
- [4] S. Ristova et al., “Energy Footprint of Microservice Topologies,” *IEEE Transactions on Software Engineering*, vol. 49, no. 4, pp. 2100–2115, 2023.
- [5] AutoMQ, “Apache Kafka vs. RabbitMQ: Differences & Comparison,” *Industry Benchmark Report*, 2024.
- [6] A. Shatnawi et al., “Migrating JSON to Protobuf for Energy Efficiency,” in *Proc. of the International Conference on Software Technologies (ICSOFT)*, 2025, pp. 112–120.